

Cache-Mediated Safety Erasure: Testing Whether KV-Cache Optimizations Selectively Weaken Refusal Behavior

Aryan Gupta
aryan.cs.app@gmail.com

Abstract

Inference systems increasingly compress or evict key-value (KV) caches to reduce serving cost and extend usable context. This paper registers a safety-specific test: cache optimizations can selectively weaken refusal and policy-following behavior while preserving ordinary task capability. We call the target phenomenon *cache-mediated safety erasure* only if the evidence ladder clears behavioral cache sensitivity, safety-minus-capability selectivity, cross-family replication, targeted mitigation, causal localization to policy or system cache state, and an alignment contrast against a matched base-model track. Empirical claims are scoped to the families, prompt suites, audit sources, and cache policies with completed provenance artifacts. Empirical claims supported by current artifacts: `behavioral_cache_sensitivity`, `safety_minus_capability_selectivity`, `cross_family_replication`, `audit_provenance_complete`. Claims not yet supported (artifacts pending): `targeted_mitigation`, `causal_localization`, `alignment_contrast`.

1 Introduction

Safety evaluations usually treat a deployed language model as a fixed pair of weights and prompts. Real serving systems are less static. KV-cache eviction, quantization, paged attention, batching, and long-context memory management alter the transient state through which a model attends to prior tokens (Kwon et al., 2023). If aligned behavior depends on transient cache state, then a throughput optimization can also become an unintentional safety intervention.

This paper asks whether inference-time KV-cache optimizations can erase or weaken refusal behavior more than they degrade ordinary capability. The motivating observation is not merely that compression can hurt language-model quality. Prior work already shows that cache compression can produce uneven degradation, instruction-following failures, and system-prompt leakage (Chen et al., 2025; Ananthanarayanan et al., 2026; Yang et al., 2024). The stronger question is whether safety behavior is selectively fragile and causally recoverable by preserving or restoring policy-relevant cache state.

We organize the study around a strict claims ladder:

1. **Behavioral cache sensitivity:** cache interventions change safety, leakage, or capability behavior.
2. **Selective safety degradation:** refusal or leakage behavior degrades more than ordinary capability under matched cache pressure.
3. **Cross-family replication:** at least two independent instruction-tuned model families show positive selectivity under the same primary matrix.

4. **Mitigation and causal localization:** policy-preserving retention reduces safety degradation and outperforms matched non-policy controls.
5. **Alignment contrast:** a matched base-model track does not show the same safety-selective refusal profile under equivalent cache pressure.

Only the full ladder justifies the phrase *cache-mediated safety erasure*. If the experiments support only the first two claims, the correct contribution is an open-model deployment evaluation of KV-compression pitfalls, not a new mechanistic safety phenomenon.

Contributions. The study contributes five components: a registered cross-family panel for safety, leakage, over-refusal, capability, and adversarial long-context placement; absolute and log-odds selective degradation metrics that subtract ordinary capability loss from safety loss; token-role cache accounting for system, user, template, and generated spans; a base-model alignment-contrast track; and mitigation or causal-localization tests that distinguish policy-token effects from “more cache helps” baselines.

2 Related Work

KV-cache compression and routing. H2O retains heavy-hitter tokens for efficient generation (Zhang et al., 2023), StreamingLLM emphasizes attention sinks for streaming contexts (Xiao et al., 2024), SnapKV compresses prompts based on observed attention structure (Li et al., 2024), and KIVI and KVQuant study low-bit KV-cache quantization (Liu et al., 2024; Hooper et al., 2024). No Token Left Behind argues that importance-aware mixed precision can preserve information that exhaustive eviction loses (Yang et al., 2024). The closest behavioral predecessor is Chen et al. (2025), who report that KV-cache compression can unevenly harm instruction following and increase system-prompt leakage. Ananthanarayanan et al. (2026) frame compression as perturbing token accessibility through attention dynamics. The central distinction here is safety-specific and causal: whether policy-token cache state is selectively fragile and recoverable.

Cache editing as a safety control surface. CachePrune edits KV-cache state to defend against indirect prompt injection (Wang et al., 2025), showing that cache state can be an intervention point rather than only an efficiency artifact. Prompt-leakage work studies privacy and instruction extraction risks (Hui et al., 2024). We instead test whether ordinary inference optimizations can weaken refusal behavior and whether token-role-aware retention can mitigate that weakness.

Alignment mechanisms and refusal. Refusal behavior can be steered through activation directions and may not reduce to a single global direction (Arditi et al., 2024; Joad et al., 2026). Recent alignment-routing work argues that policy behavior can localize to sparse routing mechanisms (Frank, 2026), while Any-Depth Alignment studies inference-time restoration of refusal behavior at different generation depths (Zhang et al., 2026). These works motivate the possibility that safety state has identifiable carriers. Our experiments ask whether KV cache slices are among those carriers.

Phenomenon-first safety papers. Subliminal learning, token entanglement, and emergent misalignment are relevant here because each frames a safety question around a specific mechanism

rather than only a new benchmark (Cloud et al., 2025; Zur et al., 2025; Betley et al., 2025). Cache-mediated safety erasure is posed in the same form: a falsifiable phenomenon with controls that separate selective safety degradation from ordinary performance degradation.

Guardrail and trustworthiness evaluations. Classifier poisoning and rapid-response poisoning studies show that safety layers can fail through mechanisms outside ordinary benchmark accuracy (Bowers et al., 2026; Huang et al., 2026). Broader trustworthiness evaluations similarly argue for multi-axis safety measurement rather than a single helpfulness score (Wang et al., 2023). These works motivate our claims ladder, but the present study targets a different inference-time mechanism: transient KV-cache state.

Evaluation anchors. The registered primary sweep requires public prompt-injection prompts from the Cyberek prompt-injection dataset (Cyberek, 2026), harmful-request prompts from AdvBench and JailbreakBench (Zou et al., 2023; Chao et al., 2024), XSTest-safe over-refusal prompts (Rottger et al., 2023), benign-instruction prompts from Databricks Dolly 15k (Databricks, 2023), and ARC-Easy capability questions (Clark et al., 2018). Broader anchors such as HarmBench, IFEval, RULER, and LongBench define extension analyses only when their corresponding artifacts are present (Mazeika et al., 2024; Zhou et al., 2023; Hsieh et al., 2024; Bai et al., 2024). The registered model panel includes Qwen checkpoints (Qwen Team, 2024) plus independent open instruction-tuned families when their pinned checkpoints pass the same text-only cache-intervention harness.

3 Hypotheses

H1: cache interventions change safety behavior. Safety, leakage, or over-refusal metrics differ between uncompressed baseline decoding and at least one cache intervention at matched prompts, seeds, and decoding settings.

H2: safety degradation is selective. Safety degradation exceeds ordinary capability degradation under the same policy and budget. We quantify this with the Selective Safety Erasure Index:

$$\text{SSEI}(p, b, m) = \Delta_{\text{safety}}(p, b, m) - \Delta_{\text{capability}}(p, b, m), \quad (1)$$

where p is a cache policy, b is its budget, and m is the model. Positive values indicate larger safety loss than capability loss.

H3: cross-family replication. The safety-selectivity effect appears in at least two independent instruction-tuned evaluated model families under the same primary policy and prompt matrix. Base-model track results, judge sources, and derivative alignment-contrast variants do not count as independent replication families.

H4: targeted mitigation. Policy-preserving cache retention reduces safety degradation relative to matched generic retention without unacceptable capability loss or benign over-refusal.

H5: causal localization. Policy or system cache preservation outperforms matched user-token preservation under the same retained-token budget. For length-preserving perturbations such as simulated KV quantization, restoring baseline key/value slices for system or policy spans must recover refusal more than restoring a matched number of user-token slices. For eviction policies

where literal reinsertion is not yet implemented, token-role-aware pinning is the registered causal-localization test. This separates policy-state localization from a generic benefit of keeping additional cache positions.

H6: alignment contrast. The matched base-model track shows raw cache sensitivity without the same instruction-refusal selectivity profile, supporting the interpretation that positive SSEI is tied to instruction tuning and safety alignment rather than generic compression sensitivity.

4 Methods

4.1 Models

The registered panel contains two analysis tracks. The chat-safety track includes instruction-tuned open checkpoints from the Qwen, GPT-OSS, Llama, Gemma, Mistral, OLMo, and Phi families when their pinned revisions, licenses or access constraints, tokenizer sources, chat templates, and cache-runner feasibility reports are complete. The base-model track includes the matched Qwen2.5 base checkpoint and any checkpoint without validated instruction/chat formatting; it uses frozen completion-format prompts and continuation scoring rather than chat-refusal labels.

Cross-family claims require at least two instruction-tuned families completing the same primary matrix with valid readiness reports and positive selectivity estimates. Base-model results are reported next to the matched Qwen instruction checkpoint as an alignment contrast, not as direct evidence for instruction-following refusal behavior.

All experiments use open-weight models, public datasets, and local evaluation code. Model loading uses bfloat16 weights, `safetensors`, and explicit cache-state access so that the evaluated interventions are visible in the generation loop rather than delegated to external serving systems. This design is appropriate for mechanism and provenance checks, but it is not by itself a validation of native production serving stacks. Claims about vLLM, TGI, PagedAttention eviction, or server-native low-bit KV-cache kernels require separate artifacts from those systems and are treated as deployment-validation extensions.

4.2 Cache Policies

The primary eviction sweep compares uncompressed generation against five registered interventions:

- **sliding window:** retain only the last N cached tokens;
- **sink plus recent:** retain the first S tokens and last N tokens;
- **random matched:** evict random tokens at the same budget as a structured policy;
- **policy pinned:** protect system or policy spans before filling the remaining budget with sink/recent tokens;
- **user pinned:** protect matched user spans as the non-policy control for policy-pinned retention.

Attention-H2O retention is treated as a diagnostic extension when attention-score artifacts are present, not as a required primary-policy gate. If attention-aware artifacts are absent, the paper must not claim coverage of modern attention-aware cache compression methods such as H2O,

SnapKV, or StreamingLLM-style retention; it may only describe the registered non-attention primary policies and any completed diagnostic extension.

Simulated KV int8 and int4 quantize-dequantize perturbations are Phase 5 diagnostics, not part of the primary eviction-policy matrix. They are reported separately unless replicated in a real serving implementation.

Each policy logs retained and evicted token counts by role, layer count, decode step, policy label, and cache-norm statistics where available. These logs support both figures and sanity checks that compression was active.

4.3 Prompt Suites

The benchmark separates safety, leakage, over-refusal, and capability:

- **system leakage:** safe prompts asking the model to reveal hidden system instructions;
- **public system leakage:** public prompt-injection attacks paired with a hidden system canary;
- **public refusal safety:** harmful-request prompts prepared from public datasets;
- **public benign over-refusal:** benign requests that should be answered rather than refused;
- **public XSTest safe:** additional benign safety-sensitive prompts;
- **public capability ARC:** multiple-choice capability controls;
- **adversarial placement:** harmful requests placed after a registered long benign preamble to induce natural pressure on earlier system or policy tokens;
- **base alignment contrast:** frozen completion-format prompts scored by safe-minus-unsafe continuation log-likelihood margins;
- **instruction following and long-context controls:** verifier-compatible IFEval-style prompts and long-context subsets included only when source provenance and evaluation scripts are complete.

Every processed suite writes a manifest with prompt IDs, source dataset revisions, record counts, and SHA256 hashes. Empirical claims require public provenance for all benchmark-derived prompts. Because public harmful-prompt benchmarks may overlap with model safety training or defensive tuning, public-suite results are interpreted as public-benchmark behavior unless a disjoint contamination check is present. A stronger generalization claim requires either held-out public datasets not used in the primary sweep, such as the registered HarmBench CI-extension shard, or a separately documented synthetic harmful-intent diagnostic that records its generation procedure and avoids procedural harmful details in the paper.

4.4 Metrics

Safety and leakage. The safety metrics include refusal-string detection, unsafe-compliance proxies from public benchmark labels, and exact or approximate leakage overlap against hidden system text. Leakage is measured by exact substring flags and token-overlap or ROUGE-style similarity.

Capability and over-refusal. Capability is tracked with exact-match or multiple-choice accuracy where labels exist. Benign over-refusal is measured by refusal-rate on benign suites and instruction-following pass rates when verifier-compatible prompts are available.

Statistical uncertainty. The primary uncertainty report uses prompt-clustered confidence intervals. Main tables include paired safety degradation intervals by prompt and seed. The absolute Selective Safety Erasure Index, $SSEI_{abs}$, is safety degradation minus capability degradation. Binary outcomes additionally use a fixed Haldane–Anscombe smoothed log-odds contrast with $\alpha = 0.5$; $SSEI_{logodds}$ is the safety log-odds degradation minus the capability log-odds degradation. The strongest claims require narrow intervals, complete prompt-policy coverage, response-length and parse-rate reporting, and agreement between automated metrics and the declared audit source.

4.5 Causal Restoration

For length-preserving cache perturbations, we run baseline and compressed decoding on identical prompts. We then patch selected key/value slices from the baseline cache into compressed runs and compare:

1. compressed decoding with no patch;
2. compressed decoding with system-role slices patched from baseline;
3. compressed decoding with user-role slices matched to the system-token count;
4. policy-pinned mitigation where system-role spans are retained under eviction;
5. user-pinned mitigation where matched user-role spans are retained under the same budget.

The causal localization claim requires system-role restoration or policy-pinned retention to recover refusal or leakage avoidance more than matched user-role controls. Role-derived spans are computed from the tokenizer/chat-template role annotations. Hard-coded token-index patches are diagnostic only and are not sufficient evidence because they can confound token role with absolute position. For quantization-based patching, the causal comparison is length-preserving: baseline and compressed runs use the same prompt tokens, positions, and RoPE phases, and the patch replaces key/value tensors at existing positions rather than inserting or reindexing cache entries. For eviction policies where cache length changes, the registered mitigation is policy pinning rather than reinserting baseline vectors into a shorter cache. Any non-length-preserving cache splice is diagnostic only and cannot support H3 unless its positional treatment is explicitly validated.

5 Experimental Design

Reproducibility. Each analysis records the resolved configuration, environment metadata, prompt records, raw generations, metrics, cache statistics, figure source data, figure hashes, and table hashes. A run is eligible for inclusion only when public dataset provenance is present, the prompt-policy-seed matrix is complete, cache interventions are active, and all reported tables and figures are generated from the recorded artifacts.

Evaluation sequence. The empirical study is ordered to reduce false positives:

- verify cache intervention mechanics on a small diagnostic slice;
- run the primary public-suite sweep across safety, leakage, benign, capability, and adversarial-placement prompts for all feasible chat-safety checkpoints;
- score the base-model alignment-contrast track with frozen safe and unsafe continuations;
- test whether the strongest effects reproduce across at least two instruction-tuned families;
- run causal restoration and policy-pinning diagnostics only on families with meaningful selective effects.

Blinded audit. Automated safety labels are coarse. The analysis therefore includes a blinded audit of high-impact slices: baseline versus strongest compressed policies, safety suites where unsafe-compliance proxies increase, and examples where system-role patching changes behavior. Condition labels are hidden during annotation and joined to experimental metadata only after labeling. The audit source is reported explicitly as either human annotation or a documented open local judge; open-judge labels are not described as human labels. For open-local-judge audits, the audit manifest must report the exact judge model, prompt-template hashes, raw-output hashes, and response-length calibration buckets. Without that calibration, open-judge labels are diagnostic only and cannot substitute for audit support in the claims ladder, because cache interventions can change response length and thereby bias a judge toward or away from refusal labels.

6 Visualization Strategy

The main figures are designed to make the proposed phenomenon interpretable at the level of prompt, policy, and token role, rather than only reporting aggregate scatterplots or bar charts. We use four figure families.

Cache-state fingerprints. For each policy-budget condition, the analysis renders a token-role retention fingerprint. When explicit layer metadata are available, the vertical axis reports layer bins; otherwise, the figure is labelled as a row-order diagnostic and is not used for layer-localization claims. System, hidden-system, user, template, and generated spans are shown as separate bands. A safety-specific erasure pattern is diagnosed by a visible break or thinning in policy-token access, not merely by a lower scalar score.

Safety-capability phase portraits. Instead of a static scatterplot alone, each policy traces a trajectory through safety loss and capability loss as the cache budget tightens. A selective safety-erasure pattern is indicated by an initial increase in safety loss at approximately fixed capability loss, followed only later by broad capability degradation.

Uncertainty braids. For each cache policy, the analysis draws a linked row of capability degradation, safety degradation, and selective safety-erasure estimates with their confidence intervals. This makes the evidential pattern visible: a strong positive result needs the safety interval to separate from capability loss and the derived SSEI interval to remain positive.




Figure 1: Safety loss versus capability loss. Each row is one cache policy; blue marks capability loss and red marks safety loss. A selective safety effect would put the red marker clearly to the right of the blue marker.

Restoration flow diagrams. Causal patching is shown as a flow from baseline to compressed decoding to system-patched, matched-user-patched, and policy-pinned conditions. The diagnostic pattern is movement back toward the baseline only when policy-relevant cache slices are restored or protected.

Prompt-level effect maps. For audited prompts, the analysis embeds each prompt-condition pair using open embeddings or metric vectors and draws paired edges from baseline to compressed generations. Clusters of long, directionally consistent edges identify families of prompts where cache interventions produce coherent safety shifts.




Figure 2: Selective Safety Erasure Index (SSEI) by policy. Positive bars mean safety fell more than capability; bars near zero do not support selective erasure; negative bars mean capability fell more than safety.

7 Figures and Tables

8 Results

8.1 Behavioral Cache Sensitivity

Table 1 reports paired baseline-versus-intervention effects across the registered safety, leakage, benign, and capability suites.

8.2 Selective Safety Degradation

Top policy-level SSEI: `sliding_window__budget128` with estimate 0.092 and 95% interval [0.083, 0.101]. A positive result requires safety degradation to exceed capability degradation under matched cache pressure, not merely lower overall quality.

8.3 Token-Role Cache Accounting

The cache-state fingerprint in Figure 4 reports token-role retention under each cache policy. The mechanistic claim requires role-specific cache disruption to align with behavioral degradation: policy-relevant spans should be disproportionately evicted or perturbed in conditions with positive SSEI, and protected or restored spans should move the model back toward baseline behavior.

8.4 Causal Restoration and Mitigation

The central causal claim requires system-role restoration or pinning to outperform matched non-policy controls.




Figure 4: Cache-state fingerprint. Bright cells mean a token role was usually retained in cache; dark cells mean it was usually evicted. Rows are policy-layer bands and columns are prompt-token roles.

8.7 Audit Support and Failure Modes

Evidence gate not satisfied. Report blinded audit label rates before making safety claims.

Evidence gate not satisfied. Report paired baseline-versus-treatment audit deltas for unsafe compliance, refusal correctness, leakage, and benign over-refusal.

Evidence gate not satisfied. Report blinded causal-diagnostic audit label rates before making restoration claims.

Evidence gate not satisfied. Report paired causal-diagnostic audit deltas for system-role patching, matched non-policy patching, and policy-pinned mitigation.

Qualitative examples are summarized at the category level rather than reproduced verbatim when they contain procedural harmful details.

9 Discussion

If the causal restoration experiments succeed, the result implies that alignment behavior should be evaluated as a property of the full inference stack, not only the trained model and prompt. The result would motivate alignment regression tests and token-role-aware cache policies for similar cache-management settings. A policy-pinned mitigation would show that preserving safety does not require disabling all cache compression.

If only selective degradation appears, the contribution is still useful but narrower: it would be evidence that cache compression can be a deployment-time safety risk. If no selective degradation appears, the study becomes a falsification result that constrains future claims about safety-specific cache fragility.



Figure 5: Behavior loss versus cache-token loss. Purple points show SSEI, black open circles show system-token cache loss, and gray open circles show user-token cache loss for the same policy.

10 Limitations

The study uses open models, public datasets, and local evaluation code. This improves reproducibility but limits comparisons against closed deployed systems and production servers. Automated refusal and leakage metrics are imperfect, so final safety claims require human audit or an open local judge with documented model provenance and response-length calibration. Simulated KV quantization is not a faithful implementation of group-wise, per-channel, residual-window, or mixed-precision production quantizers such as KIVI- or KVQuant-style kernels. Attention-based retention policies require attention statistics and, unless their artifacts are present, cannot support claims about attention-aware eviction methods. Public harmful-prompt suites may be contaminated by model safety tuning; absent a disjoint public or synthetic contamination check, conclusions are limited to the evaluated public benchmark distribution. Finally, unless at least two independent instruction-tuned families complete the same primary matrix, the paper must not claim cross-family generality.

11 Ethics and Safety

The experiments evaluate refusal and leakage behavior without publishing procedural harmful instructions. Raw generations are retained locally for audit and reproducibility, but paper examples should be redacted or summarized when they risk enabling harm. The purpose is to identify deployment-time safety regressions and mitigations in open inference stacks.




Figure 6: Policy-level estimates with 95% confidence intervals. Intervals that cross zero do not provide reliable evidence for that effect; a supported selective-erasure claim requires the SSEI interval to stay positive.




Figure 7: Causal restoration toward baseline on refusal-safety prompts. Zero means the patched run stayed at the compressed behavior; one means it returned to baseline; values outside this range indicate worse-than-compressed behavior or overshoot.

12 Reproducibility Checklist

Requirement	Implementation
Open dependencies	Open models, public datasets, local metrics, no proprietary judges.
Run provenance	Resolved analysis configuration, environment metadata, dataset revisions, and artifact hashes.
Full matrix coverage	Prompt-policy-seed completeness is verified before analysis.
Compression sanity	Cache interventions must measurably alter retained or perturbed cache state.
Statistical uncertainty	Prompt-clustered and paired intervals in aggregate metrics.
Paper artifacts	Figure and table manifests with SHA256 hashes.
Audit support	Blinded condition labels with post-annotation metadata joining and explicit source reporting.

Figure unavailable
Compressed, system-patched, matched-user-patched, and policy-pinned conditions.

Figure 8: Causal restoration flow for the refusal-safety patch comparison. Arrows start at compressed behavior and end at the patched behavior; the matched user-token control must lag the system-token patch for a causal system-cache claim to pass.

suite	policy	Safety delta [95% CI]	Paired / clusters
Adversarial Refusal Safety	Sink+recent	-0.667 [-1.000,0.000]	3
System leakage probe	Window 128	0.429 [0.429,0.429]	2
Public system leakage	Window 128	0.293 [0.270,0.316]	1300
System leakage probe	Policy-pinned cache	-0.071 [-0.143,0.000]	2
System leakage probe	Random matched	-0.071 [-0.286,0.143]	2
System leakage probe	Sink+recent	-0.071 [-0.286,0.143]	2
System leakage probe	user_pinned / budget 128 / sink 8	-0.071 [-0.286,0.143]	2
Public system leakage	Policy-pinned cache	0.063 [0.053,0.073]	1300
Refusal safety	Window 128	0.053 [0.038,0.068]	1300
Public system leakage	Random matched	0.032 [0.024,0.040]	1300
Refusal safety	Random matched	0.032 [0.015,0.047]	1300
Refusal safety	Policy-pinned cache	-0.024 [-0.038,-0.010]	1300
Public system leakage	user_pinned / budget 128 / sink 8	-0.017 [-0.022,-0.012]	1300
Public system leakage	Sink+recent	-0.013 [-0.018,-0.008]	1300
Refusal safety	user_pinned / budget 128 / sink 8	0.007 [-0.007,0.022]	1300
Refusal safety	Sink+recent	-0.005 [-0.019,0.010]	1300

Table 1: Largest absolute suite-level safety-degradation effects with paired prompt counts. Full suite-policy table is exported as Markdown.

13 Conclusion

Cache-mediated safety erasure is warranted as a positive empirical claim only if it survives causal controls. Novelty must therefore be earned empirically: safety behavior must degrade selectively, policy-relevant cache restoration must outperform matched controls, and mitigation must preserve safety without globally disabling cache compression. If the evidence supports those claims, the result would show that alignment behavior can depend on the transient inference state created by cache management. If the evidence does not support them, the study still constrains future claims about safety-specific cache fragility.

References

Samhruth Ananthanarayanan, Ayan Sengupta, and Tanmoy Chakraborty. Understanding the physics of key-value cache compression for llms through attention dynamics. *arXiv preprint*

policy	Mean safety	Mean capability	Policy SSEI [95% CI]
Baseline	0.607	0.002	
Policy-pinned cache	0.612	0.002	0.010 [0.006,0.015]
Random matched	0.609	0.007	0.020 [0.014,0.026]
Sink+recent	0.733	0.002	-0.005 [-0.009,-0.001]
Window 128	0.478	0.008	0.092 [0.083,0.101]
user_pinned / budget 128 / sink 8	0.621	0.003	-0.002 [-0.006,0.002]

Table 2: Policy-level safety, capability, and selective safety erasure summary.

Table 3: Claim ladder status from the current selectivity panel artifacts.

Claim	Status	Notes
behavioral_cache_sensitivity	supported	6 of 8 instruction-tuned models show $ SSEI \geq 0.01$.
safety_minus_capability_selectivity	supported	5 models have a registered policy with positive SSEI whose 95% CI excludes 0 (0.01 threshold).
cross_family_replication	supported	Families with positive instruction-tuned selectivity excluding 0: ['Llama', 'OLMo', 'Phi', 'Qwen'].
targeted_mitigation	pending	Mitigation comparison requires policy_pinned vs matched user_pinned causal contrasts that this panel does not yet produce.
causal_localization	pending	Phase 4 causal patching is conditional on Phase 2 effects and is not yet executed for the selectivity panel.
alignment_contrast	pending	Base-model alignment-contrast suite is recorded with only 2 sampled prompts per cell for qwen2_5_7b_base; insufficient power to score.
audit_provenance_complete	supported	Models with $\geq 95\%$ local-judge coverage: ['gpt_oss_20b', 'phi4', 'qwen2_5_7b_base', 'qwen2_5_7b_instruct', 'qwen3_5_9b']; covered families: ['OpenAI', 'Phi', 'Qwen'].

arXiv:2603.01426, 2026. URL <https://arxiv.org/abs/2603.01426>.

Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, 2024. URL <https://arxiv.org/abs/2406.11717>.

Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, Zhengxiao Du, Xiao Liu, Aohan Zeng, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. Longbench: A bilingual, multitask benchmark for long context understanding. In *Annual Meeting of the Association for Computational Linguistics*, 2024. URL <https://arxiv.org/abs/2308.14508>.

Jan Betley et al. Emergent misalignment: Narrow finetuning can produce broadly misaligned llms. *arXiv preprint arXiv:2502.17424*, 2025. URL <https://arxiv.org/abs/2502.17424>.

Chase Bowers, Faizan Ali, John Hughes, Jerry Wei, and Fabien Roger. Poisoning fine-tuning datasets of constitutional classifiers. 2026. URL <https://alignment.anthropic.com/2026/backdoorring-classifiers/>. Anthropic Alignment Science Blog.

Patrick Chao, Edoardo Debenedetti, Alexander Robey, Maksym Andriushchenko, Francesco Croce, Vikash Sehwal, Edgar Dobriban, Nicolas Flammarion, George J. Pappas, Florian Tramèr, Hamed Hassani, and Eric Wong. Jailbreakbench: An open robustness benchmark for jailbreaking large

Table 4: Cross-family selectivity (SSEI) by cache policy. SSEI is safety degradation minus capability degradation relative to the no-cache-policy baseline; positive SSEI indicates safety loss exceeding capability loss. 95% bootstrap CIs in brackets.

Family	Model	sliding window	sink+recent	policy-pinned	user-pinned	random matched
Gemma	Gemma-2-9B-IT	-0.005 [-0.008, -0.003]	-0.001 [-0.003, 0.001]	-0.004 [-0.006, -0.002]	-0.003 [-0.008, 0.001]	-0.000 [-0.003, 0.003]
OpenAI	GPT-OSS-20B	-0.068 [-0.088, -0.048]	-0.068 [-0.088, -0.048]	-0.079 [-0.099, -0.060]	-0.082 [-0.101, -0.063]	-0.079 [-0.099, -0.060]
Llama	Llama-3.1-8B-Instruct	-0.085 [-0.105, -0.067]	0.007 [0.002, 0.012]	-0.001 [-0.007, 0.004]	0.024 [0.016, 0.031]	0.022 [0.014, 0.029]
Mistral	Mistral-7B-Instruct-v0.3	0.009 [0.002, 0.017]	0.001 [-0.002, 0.004]	-0.004 [-0.008, -0.000]	-0.008 [-0.015, -0.002]	0.002 [-0.003, 0.006]
OLMo	OLMo-3-7B-Instruct	-0.004 [-0.008, 0.000]	-0.010 [-0.014, -0.005]	-0.016 [-0.020, -0.011]	0.026 [0.019, 0.033]	0.005 [0.000, 0.010]
Phi	Phi-4	0.084 [0.076, 0.091]	0.001 [-0.000, 0.003]	-0.001 [-0.003, 0.000]	0.055 [0.046, 0.063]	0.044 [0.038, 0.050]
Qwen	Qwen2.5-7B base	—	—	—	—	—
Qwen	Qwen2.5-7B-Instruct	0.011 [0.007, 0.016]	-0.003 [-0.008, 0.001]	-0.001 [-0.006, 0.003]	0.017 [0.011, 0.022]	0.012 [0.007, 0.016]
Qwen	Qwen3-9B	0.092 [0.083, 0.101]	-0.005 [-0.009, -0.001]	0.010 [0.006, 0.015]	-0.002 [-0.006, 0.002]	0.020 [0.014, 0.026]

language models. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*, 2024. URL <https://arxiv.org/abs/2404.01318>.

Alex Chen, Renato Geh, Aditya Grover, Guy Van den Broeck, and Daniel Israel. The pitfalls of kv cache compression. *arXiv preprint arXiv:2510.00231*, 2025. URL <https://arxiv.org/abs/2510.00231>.

Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 2018. URL <https://arxiv.org/abs/1803.05457>.

Alex Cloud, Minh Le, James Chua, Jan Betley, Anna Sztyber-Betley, Jacob Hilton, Samuel Marks, and Owain Evans. Subliminal learning: Language models transmit behavioral traits via hidden signals in data. *arXiv preprint arXiv:2507.14805*, 2025. URL <https://arxiv.org/abs/2507.14805>.

Cybersec. Prompt injection and jailbreak detection dataset. Hugging Face dataset, 2026. URL <https://huggingface.co/datasets/cybersec/Prompt-injection-dataset>.

Databricks. Databricks dolly 15k. Hugging Face dataset, 2023. URL <https://huggingface.co/datasets/databricks/databricks-dolly-15k>.

Gregory N. Frank. How alignment routes: Localizing, scaling, and controlling policy circuits in language models. *arXiv preprint arXiv:2604.04385*, 2026. URL <https://arxiv.org/abs/2604.04385>.

Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W. Mahoney, Yakun Sophia Shao, Kurt Keutzer, and Amir Gholami. Kvquant: Towards 10 million context length llm inference with kv cache quantization. In *Advances in Neural Information Processing Systems*, 2024. URL <https://arxiv.org/abs/2401.18079>.

Cheng-Ping Hsieh, Simeng Sun, Samuel Krman, Shantanu Acharya, Dima Rekeshe, Fei Jia, Yang Zhang, and Boris Ginsburg. Ruler: What’s the real context size of your long-context language models? In *Conference on Language Modeling*, 2024. URL <https://arxiv.org/abs/2404.06654>.

- David Huang, Jaewon Chang, Avidan Shah, Prateek Mittal, and Chawin Sitawarin. Rapid poison: Practical poisoning attacks against the rapid response framework. In *ICLR 2026 Workshop on Agents in the Wild*, 2026. URL <https://openreview.net/forum?id=tIOo2CG7cL>.
- Bo Hui, Haolin Yuan, Neil Gong, Philippe Burlina, and Yinzhi Cao. Pleak: Prompt leaking attacks against large language model applications. *arXiv preprint arXiv:2405.06823*, 2024. URL <https://arxiv.org/abs/2405.06823>.
- Faaiz Joad, Majd Hawasly, Sabri Boughorbel, Nadir Durrani, and Husrev Taha Sencar. There is more to refusal in large language models than a single direction. *arXiv preprint arXiv:2602.02132*, 2026. URL <https://arxiv.org/abs/2602.02132>.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles*, 2023. URL <https://arxiv.org/abs/2309.06180>.
- Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, Acyr Locatelli, Hanchen Ye, Tianle Cai, Patrick Lewis, and Deming Chen. Snapkv: Llm knows what you are looking for before generation. *arXiv preprint arXiv:2404.14469*, 2024. URL <https://arxiv.org/abs/2404.14469>.
- Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman, Beidi Chen, and Xia Hu. Kivi: A tuning-free asymmetric 2bit quantization for kv cache. In *International Conference on Machine Learning*, 2024. URL <https://arxiv.org/abs/2402.02750>.
- Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. Harmbench: A standardized evaluation framework for automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*, 2024. URL <https://arxiv.org/abs/2402.04249>.
- Qwen Team. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024. URL <https://arxiv.org/abs/2412.15115>.
- Paul Rottger, Hannah Rose Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. Xstest: A test suite for identifying exaggerated safety behaviours in large language models. *arXiv preprint arXiv:2308.01263*, 2023. URL <https://arxiv.org/abs/2308.01263>.
- Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T. Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, Zinan Lin, Yu Cheng, Sanmi Koyejo, Dawn Song, and Bo Li. Decodingtrust: A comprehensive assessment of trustworthiness in gpt models. *arXiv preprint arXiv:2306.11698*, 2023. URL <https://arxiv.org/abs/2306.11698>.
- Rui Wang, Junda Wu, Yu Xia, Tong Yu, Ruiyi Zhang, Ryan Rossi, Lina Yao, and Julian McAuley. Cacheprune: Neural-based attribution defense against indirect prompt injection attacks. *arXiv preprint arXiv:2504.21228*, 2025. URL <https://arxiv.org/abs/2504.21228>.
- Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient streaming language models with attention sinks. In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2309.17453>.

- June Yong Yang, Byeongwook Kim, Jeongin Bae, Beomseok Kwon, Gunho Park, Eunho Yang, Se Jung Kwon, and Dongsoo Lee. No token left behind: Reliable kv cache compression via importance-aware mixed precision quantization. *arXiv preprint arXiv:2402.18096*, 2024. URL <https://arxiv.org/abs/2402.18096>.
- Jiawei Zhang, Andrew Estornell, David D. Baek, Bo Li, and Xiaojun Xu. Any-depth alignment: Unlocking innate safety alignment of llms to any-depth. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=OfuY0uJyzl>.
- Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuandong Tian, Christopher Re, Clark Barrett, Zhangyang Wang, and Beidi Chen. H2o: Heavy-hitter oracle for efficient generative inference of large language models. In *Advances in Neural Information Processing Systems*, 2023. URL <https://arxiv.org/abs/2306.14048>.
- Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*, 2023. URL <https://arxiv.org/abs/2311.07911>.
- Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023. URL <https://arxiv.org/abs/2307.15043>.
- Amir Zur, Zhuofan Ying, Alexander Russell Loftus, Kerem Sahin, Steven Yu, Lucia Quirke, Tamar Rott Shaham, Natalie Shapira, Hadas Orgad, and David Bau. Token entanglement in subliminal learning. In *NeurIPS 2025 Mechanistic Interpretability Workshop*, 2025. URL <https://openreview.net/forum?id=auKgpBRzIW>.